Experiences with AWS EKS

Kubernetes and Cloud Native Group Dresden

ALFRED KROHMER

LogMeIn®
Be Limitless.

Structure

- Introduction
  - overview / EKS cluster creation
  - reasons for our ongoing migration (and what we had before)
- AWS-specific cluster addons
  - VPC CNI plugin
  - external-dns
  - alb-ingress-controller
  - kube2iam
- Worker node management
  - cluster-autoscaler
- Q&A

Introduction

- **Overview**
  - EKS = Elastic Kubernetes Service
  - AWS provisions and maintains master nodes
  - you provision and maintain worker nodes
    - “EKS-optimized” AMI by AWS
  - authentication via »AWS IAM Authenticator for Kubernetes«

    https://github.com/kubernetes-sigs/aws-iam-authenticator

Introduction

- **Cluster overview**

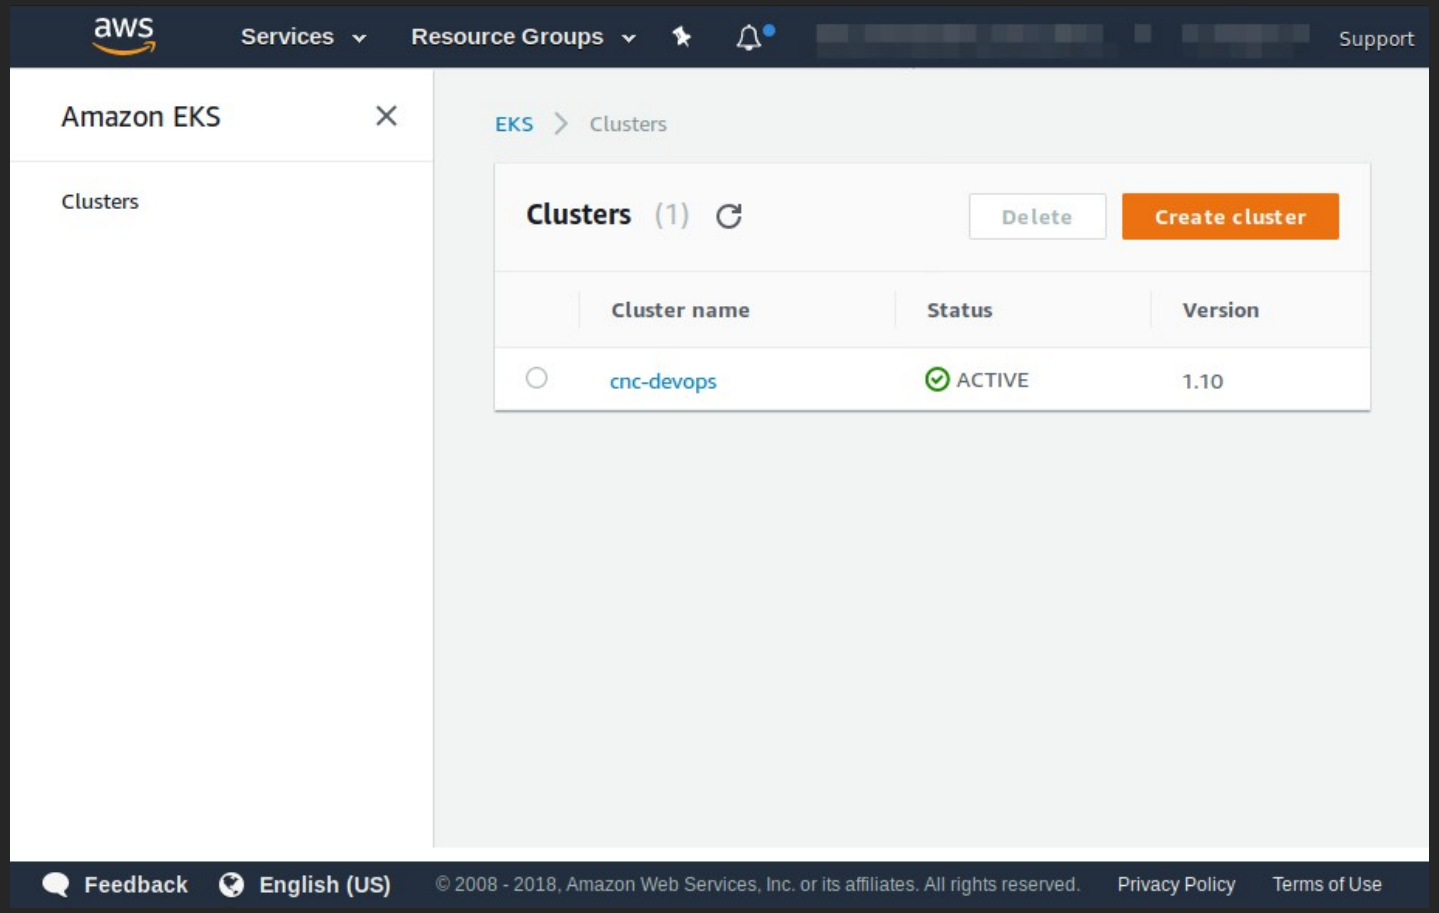

Introduction

- **Cluster creation**

  IAM user which creates the cluster is (and stays) Kubernetes admin user (for now)

  → use a headless account to create the cluster

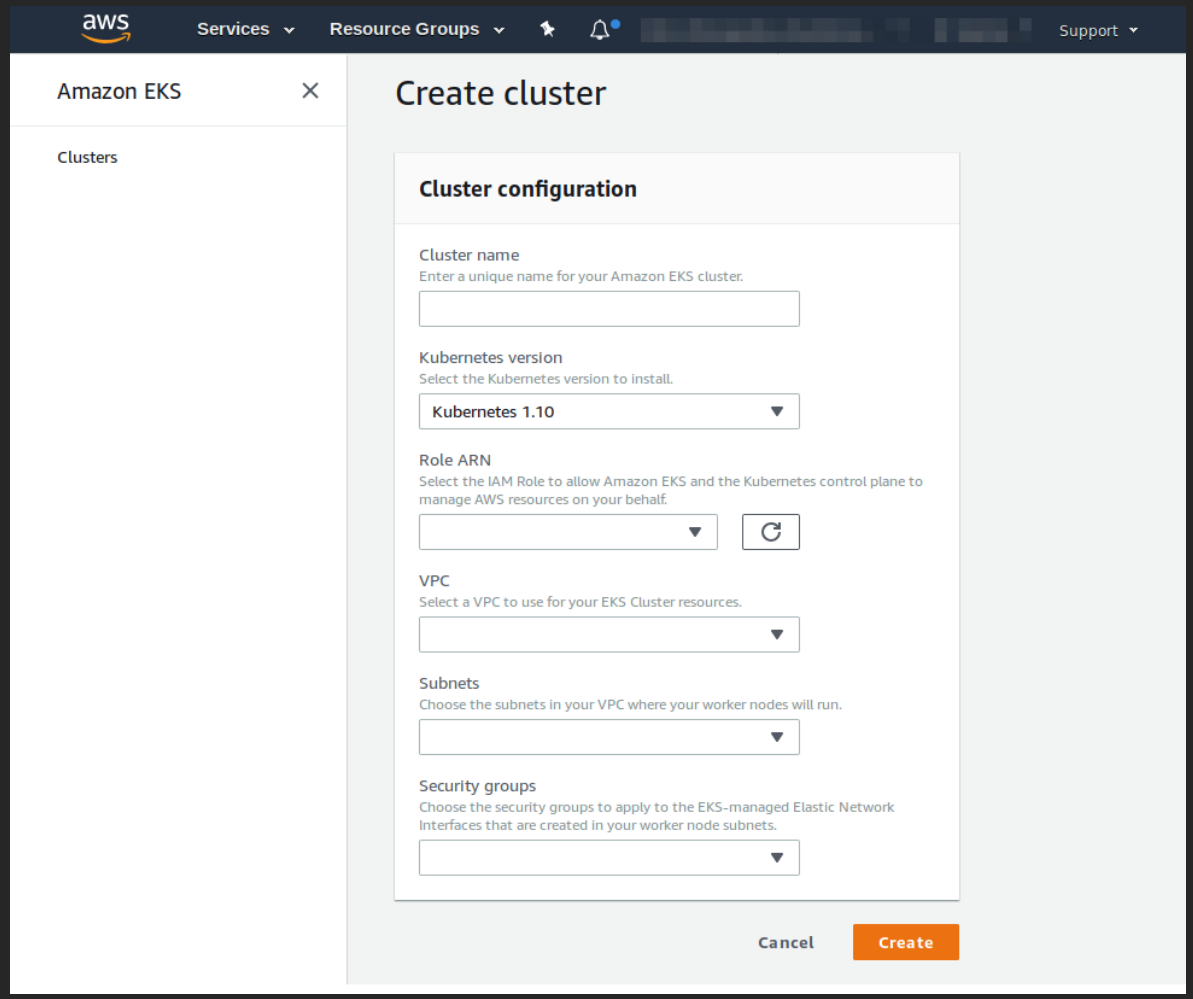

Introduction

- **Cluster details**

  → details cannot be modified (for now)

  → “Platform versions” are supposed to be rolled out by AWS

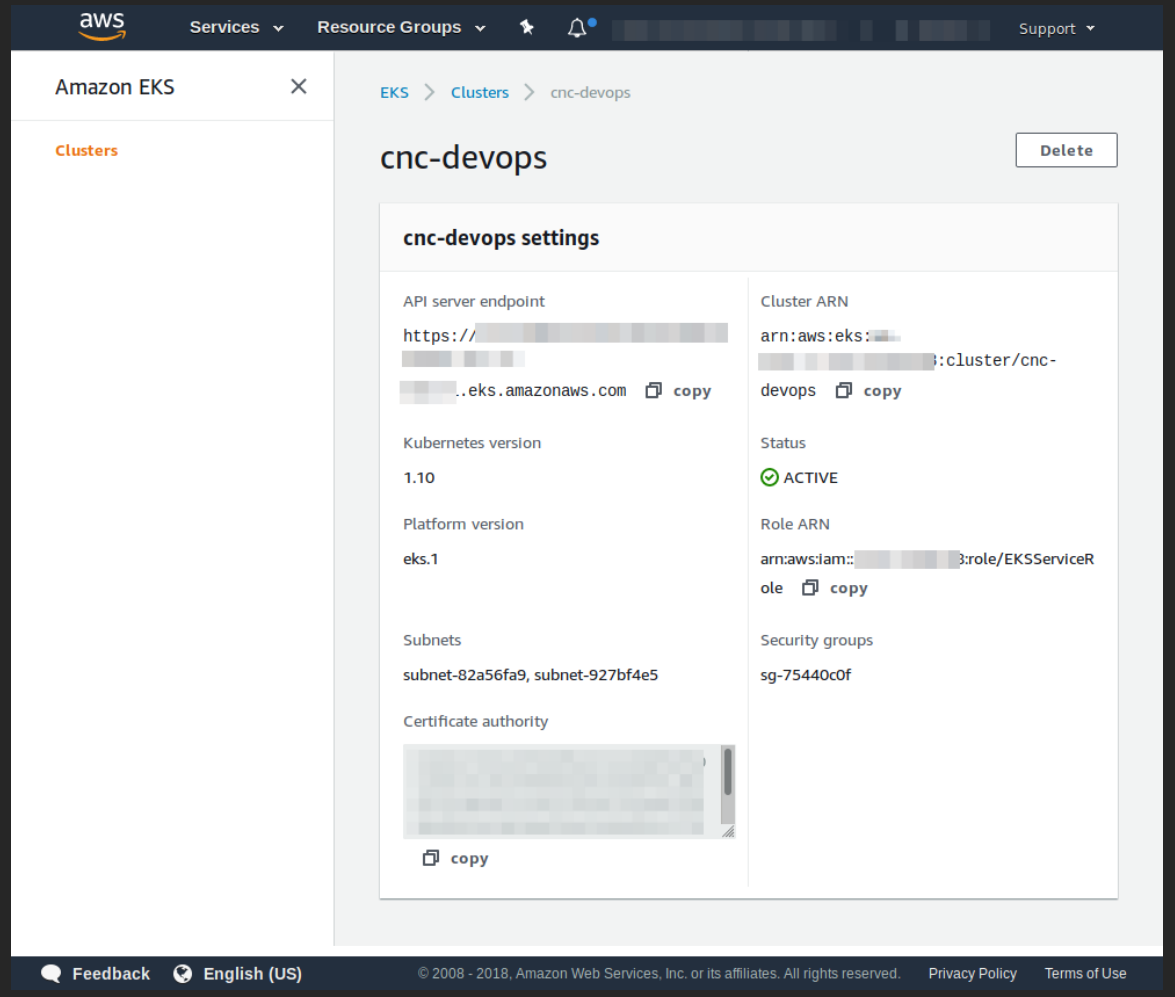

Introduction

- **Our (ongoing) migration away from**
  - **AWS ECS:**
    - slow (especially with CloudFormation)
    - lack of flexibility
    - vendor lock-in
    - EC2 instances unmaintained
    - Fargate much more expensive then EC2 spot
    - 30s maximum graceful termination period
    - cluster needs to have enough capacity **before** scheduling a task / service

Introduction

- **Our (ongoing) migration away from**
  - **our own container orchestration framework**
    - availability problems (in dev)
    - lack of flexibility
      - one “pod” per host deployments only
    - overly complicated setup with a custom host agent and in-container agent both connecting to ZooKeeper, with frontend UI and REST API
      - starting containers
      - providing application configuration

Introduction

- Our (ongoing) migration away from
  - our own container orchestration framework

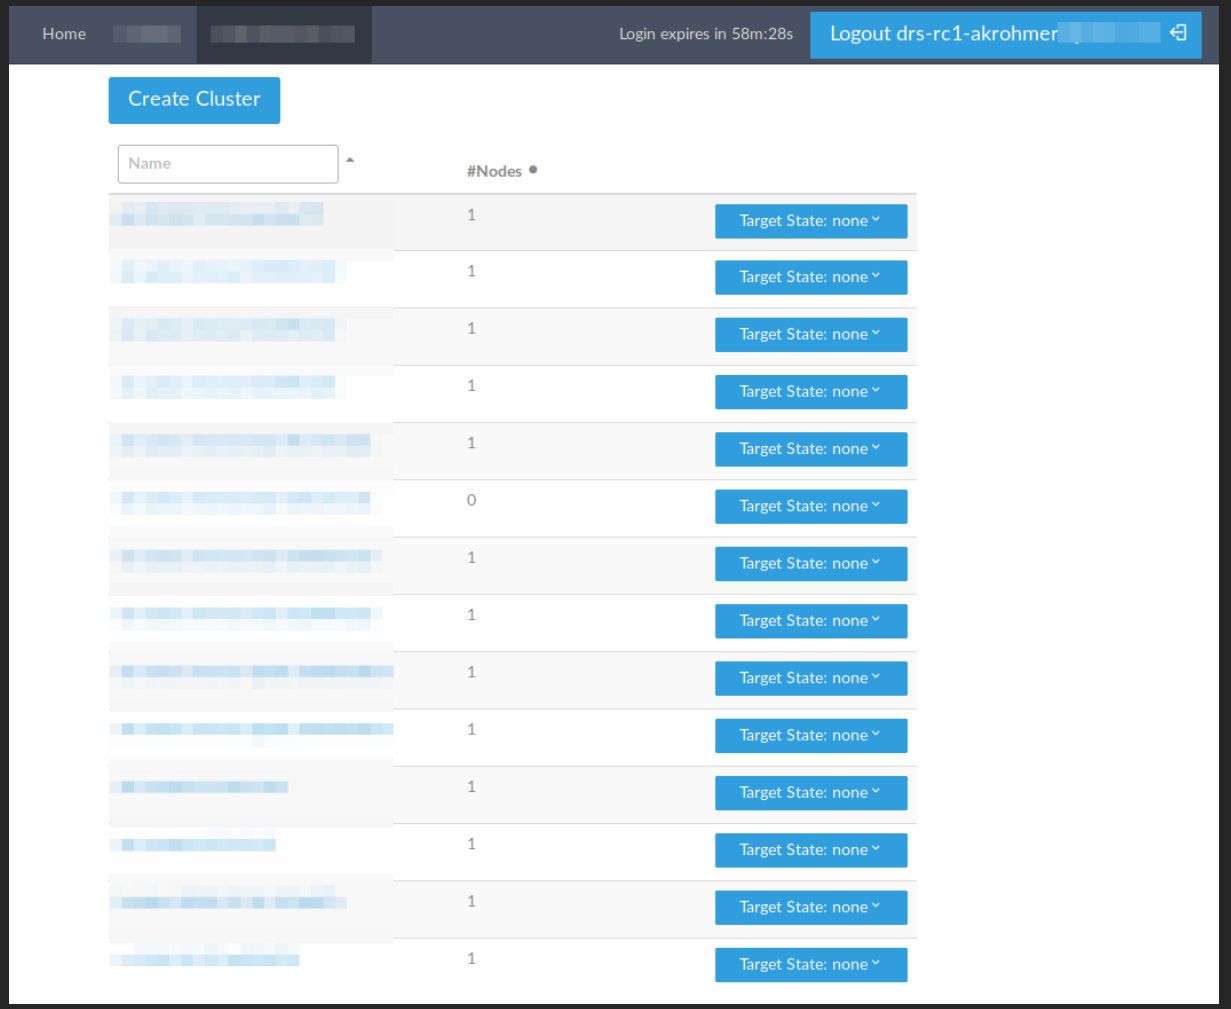

Introduction

- **Our (ongoing) migration away from**
  - **our own container orchestration framework**

  all nodes hosted via EC2

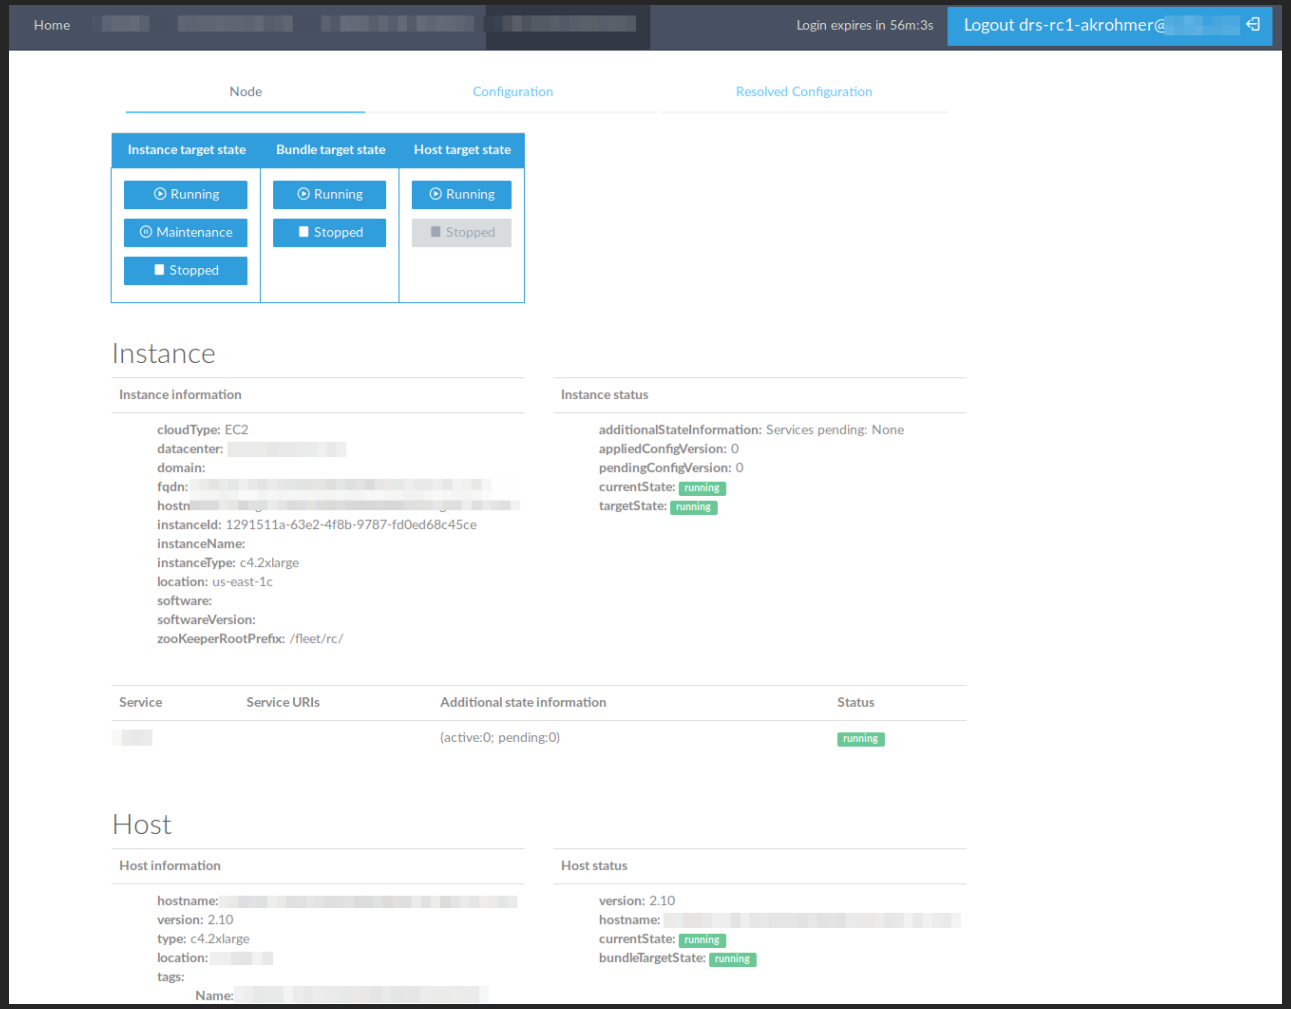

AWS-specific cluster addons

- **VPC CNI plugin**
  - CNI plugin for kubelet
  - open source
  - can be used without EKS / with custom Kubernetes installations in EC2

  - https://github.com/aws/amazon-vpc-cni-k8s

AWS-specific cluster addons

- **VPC CNI plugin**
  - uses AWS VPC as “overlay” network for Kubernetes
    - assigns secondary IPs to EC2 worker nodes
    - routing rules pass traffic through veth pairs to pod network
  - all other EC2 instances in your VPC can connect to your pods directly

AWS-specific cluster addons

- **VPC CNI plugin**
  - beware of the defaults:
    - allocates as many IPs as possible per instance → IP starvation in small subnets
    - secondary network interfaces: uses SNAT to primary private IP to reach outside of VPC
  - “number of available IPs” is not managed as a resource

AWS-specific cluster addons

- **external-dns**
  - creates Route 53 resource records for your services, ingresses and StatefulSet (or single) pods

```yaml
apiVersion: v1
kind: Service
metadata:
  name: nginx
  annotations:
    external-dns.alpha.kubernetes.io/hostname: nginx.external-dns-test.my-org.com.
spec:
  type: LoadBalancer
  ports:
  - port: 80
    name: http
    targetPort: 80
  selector:
    app: nginx
```

```yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: foo
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - backend:
          serviceName: foo
          servicePort: 80
```

  - https://github.com/kubernetes-incubator/external-dns/blob/master/docs/tutorials/aws.md

AWS-specific cluster addons

- **external-dns**
  - uses Route 53 TXT records for own storage
  - also supports other cloud providers
  - missing feature for us:
    creating per-pod DNS records

AWS-specific cluster addons

- **alb-ingress-controller**
  - creates AWS Application Load Balancers (ALBs) for your ingresses

    targets are the VPC IPs of your pods

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: foo
  annotations:
    kubernetes.io/ingress.class: "alb"
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - backend:
          serviceName: foo
          servicePort: 80
```

Listener Rule

Target group

  - https://github.com/kubernetes-sigs/aws-alb-ingress-controller

AWS-specific cluster addons

- **alb-ingress-controller**
  - stops “controlling” an ALB if you add non-existing services
    → can bring down other services behind the same ALB

AWS-specific cluster addons

- **kube2iam**
  - deployed as daemon set
  - intercepts traffic to EC2 metadata service at http://169.254.169.254/ via iptables and serves its own metadata service for pods
  - https://github.com/jtblin/kube2iam

```
apiVersion: v1
kind: Pod
metadata:
  name: aws-cli
  labels:
    name: aws-cli
  annotations:
    iam.amazonaws.com/role: role-arn
spec:
  containers:
  - image: fstab/aws-cli
    command:
      - "/home/aws/aws/env/bin/aws"
      - "s3"
      - "ls"
      - "some-bucket"
    name: aws-cli
```

AWS-specific cluster addons

- **kube2iam**
  - allows to assign IAM roles to your pods
  - whitelist allowed pod roles per namespace via annotation
  - EC2 instance role needs to be allowed to assume all pod roles

Worker node management

- **provisioning auto scaling groups (ASGs)**
  - via CloudFormation
  - one stack per ASG
  - one ASG per
    - instance type
    - launch type (on-demand / spot)
    - availability zone
    - service type (RTC / non-RTC, special kernel settings)

    → huge matrix of ASGs

Worker node management

- **provisioning auto scaling groups (ASGs)**
  - description of the matrix of all required combinations in YAML file
  - read and excuted by python script
  - still missing:
    - automated roll-over of instances
    - automated deletion of unneeded ASGs

Worker node management

- **cluster-autoscaler**
  - another cluster-addon
  - supports multiple cloud providers
  - scales ASGs up & down according to number of pending pods and underutilized nodes
  - can be installed per-namespace to support fine-granular scaling behavior

  - https://github.com/kubernetes/autoscaler/tree/master/cluster-autoscaler

Worker node management

- **cluster-autoscaler**
  - beware of the defaults:
    - stops working for some time if some actions fail
    - can scale up all ASGs to their maximum number of instances
  - don’t deploy it on spot instances!


Q&A

Thank you.